

Stage 0: Adversarial attacks and anomaly detection in MIR

Vinod Subramanian, QMUL, ROLI Ltd.

Abstract

The abstract goes here.

Index Terms

MIR

I. INTRODUCTION

Deep learning is having an increasing presence in the day to day lives of people around the world. In recent years with the improving computing capabilities of devices the public is utilising deep learning algorithms in their day to day lives. Whether it be voice recognition systems, music recommendation algorithms, tv show recommendation algorithms, filters to warp images they all use different types of deep learning approaches.

One of the reasons for the success of deep learning is the fact that it can learn complex non-linear functions with relative ease. It does this with the use of many hidden layers. According to the universal approximation theorem [1] it can approximate any function. The flipside to using deep learning models is that it is difficult to understand what a deep learning model is actually learning.

This inability to understand what is being learned by a deep learning model makes it difficult to understand whether a deep learning model will perform reliably when deployed for public use.

- 1) Introduce the increasing importance of data and how there is an abundance of large collections of data in the world
- 2) Give examples of movie collections, music streaming, personal information, instagram, youtube etc.
- 3) This data has facilitated research in deep learning. So a combination of an abundance of data and better computing has brought about a new age of deep learning
- 4) Challenge with larger datasets is mislabelling, data corruption, noise which are currently being done by mostly manual curation there is a need to work on automatic methods to curate. Practically this manifests as when you have live data coming in quickly identifying an error.
- 5) Challenge with deep learning is that while deploying them in the real world you need to be sure they are working well and always perform as expected, in effect they must be robust to noise
- 6) ROLI is interested in problems such as acoustic instrument dataset creation where high quality recordings are key, so an automatic method to identify faulty recordings that don't meet the standard
- 7) ROLI is interested in high accuracy and highly robust deep learning models that they can confidently release to their customers so there is a need to do research in investigating robustness

II. BACKGROUND

I'd like to see if I can create a summary of some of the research techniques used in some of the literature I am reviewing in the form of a table to compare techniques and performance.

A. What are adversarial attacks?

Szegedy et al. [2] discovered that in object recognition tasks by applying an imperceptible non-random perturbation to the input image the output of the network can be changed. The term "Adversarial examples" is used to describe these perturbed examples. These adversarial examples were attributed to the fact that there are blindspots in the training process for these deep learning models.

Goodfellow, Shlens and Szegedy [2] challenged the idea that adversarial examples were due to blind spots in high dimensional spaces. Instead, they suggested that they are caused due to linearity in deep learning models. According to them LSTMs, ReLUs, maxout networks, CNNs etc. were intentionally designed to behave more linearly in order to make them easier to optimize.

Assume we have a linear classifier defined by the relationship $y = f(x)$ where $f(x)$ is given by $f(x) = \omega^T x$. Our goal is to perturb the input x with perturbation η subject to the condition that $D(x + \eta, x) < \epsilon$ where D is some distance measure. We can write the new equation for the output of the classifier as:

$$\tilde{y} = \omega^T x + \omega^T \eta \quad (1)$$

By choosing η carefully so that it is aligned in the direction of the weights of the classifier ω the change in the output can be maximized.

Goodfellow, Shlens and Szegedy [?] used this assumption on deep learning models to generate adversarial examples and came up with a family of fast gradient approaches to generate adversarial examples. The success of these adversarial examples lent credence to the fact that deep learning models are susceptible to adversarial examples on account of being too linear.

Before we give more details about different types of adversarial attacks and defenses against these attacks we'll list the different categories of adversarial attacks. The category of an adversarial attack is determined by either the goal of the adversary or the information the adversary has of the classifier.

Goal of the adversary:

- 1) Untargeted attack - The adversary's goal is to simply misclassify the input. As long as the new target class of the classifier is different from the original the adversary has achieved its goal.
- 2) Targeted attack - The adversary's goal is to change the label from the original to a different label that is specified before the attack begins. Because there are more constraints in the targeted attack it is typically harder to generate targeted adversarial examples as opposed to untargeted adversarial examples.

Information of the adversary:

- 1) Perfect knowledge - In this scenario the adversary has perfect knowledge of the classifier such as the feature space, the weights of the model and the type of classifier. This is also known as a white box attack.
- 2) Limited knowledge - The classifier does not know the training data or the trained model but has knowledge about the feature representation and the type of classifier. This could be considered a black box attack.
- 3) Zero knowledge - This is a special case where an adversarial example is generated for one classifier and then tested on a different classifier which it has no information about.

In the following sections I will introduce different techniques for attacks, different techniques for defenses and some unique properties of adversarial attacks such as universal perturbation, transferability of adversarial examples and adversarial examples in the real world. Finally I will highlight the existing body of work that explores adversarial attacks in audio and talk about existing deep learning tasks in music information retrieval (MIR) that would be a good starting point for research into adversarial attacks in music.

B. Summary of attacks

C. Summary of defenses

D. Adversarial attacks in audio

E. Adversarial attacks

Deep learning has gained a lot of popularity in recent years with success in image recognition, text-to-speech, motion tracking and singing voice transcription. The success of deep learning is attributed to the fact that large non-linear models can be trained using back propagation. The tradeoff with these complex, non-linear models is that it is difficult to interpret what is being learned by the deep learning model.

Following this research Goodfellow, Shlens and Szegedy [?] suggested that adversarial examples exist due to the compounding effect of dot products through each layer of the neural network. This would mean that deep learning models are too linear. This idea can be expressed by the following equation:

If the perturbation η is aligned with the weights then the perturbation has the potential to change the output significantly. The argument was that a lot of LSTMs, ReLUs etc. were intentionally designed to behave linearly in order to make them easier to optimize. So, if we assume linearity, by using the formula listed above we can use the sign of the weights to analytically generate perturbations easily. This method for generating adversarial examples was called "fast gradient sign method".

Biggio et al. [?] approached the problem of adversarial attacks from the perspective of security and defined broadly two scenarios for attack:

- 1) Perfect knowledge:
- 2) Limited knowledge:
- 1) Discovery of adversarial examples where small perturbations in data causes high accuracy deep learning models to misclassify the input
- 2) Further, these adversarial examples were transferable between different models which is alarming
- 3) More and more efficient and cheap ways to generate adversarial attacks were discovered in an attempt to understand the underlying problem of deep learning models
- 4) Simultaneous work on solving these attacks became popular with distilled learning methods
- 5) More state of the art approaches towards defending from attacks happened
- 6) Most research was in computer vision or natural language processing, Carlini found that some of these adversarial attacks work in audio as well
- 7) Bob Sturm showed that these attacks exist in music too

F. Anomaly detection

- 1) Summarize the survey paper on anomaly detection from 2007. This survey paper talks a lot about older approaches to anomaly detection that are still relevant
- 2) For mechanical failures there are deep learning approaches for anomaly detection, mostly inspired by the fact that it is harder to define the anomaly.
- 3) Work done on anomaly detection in medicine to identify diseases from measurements automatically
- 4) Work done on anomaly detection on GTZAN dataset using classical machine learning and on mammal sound recognition

III. RESEARCH QUESTIONS AND METHODS

In this section I'll list all of my ideas for this project so far which I will reference again in the timeline section. I am going to try and include block diagrams and images to make my ideas clearer.

A. Recreating adversarial attack experiments in audio

- 1) Transferring adversarial examples in audio between models with identical inputs: In computer vision they found that adversarial examples generated in one model served as an adversarial example in another model. This indicates that there seems to be some fundamental way in which we train deep learning models that make them susceptible to errors
- 2) Can adversarial examples be played over the air: In computer vision a photo of an adversarial example taken from a phone camera still works as an adversarial example
- 3) Can you create a universal perturbation that has a high chance of making an audio an adversarial example: In computer vision they came up with a universal method to create an adversarial example

B. Adversarial attacks for audio

- 1) Unlike computer vision audio domain uses different types of input features for deep learning. It would be interesting to see if it is possible to design an adversarial attack in the time domain that is robust to the different input representations
- 2) Design attacks for a specific task such as singing voice transcription. Singing voice transcription has a few different high accuracy models to work with so it is a good candidate to perform adversarial attacks on
- 3) Some research is done showing that mp3 compression preserves the adversarial examples but it would be interesting to try different bitrates and different compression methods

C. Defenses against attacks

- 1) Apply existing defenses to the audio problem to see if they work
- 2) Design defenses that are tailored to the specific tasks that I choose to work on
- 3) Explore the idea of using adversarial examples as a data augmentation technique to create higher accuracy models

D. Anomaly detection in note level datasets

- 1) Ensemble of methods to identify poor quality, corrupted audio and noise in datasets
- 2) Replace ensemble of methods with a single deep learning model and explore the use of multi-task learning

IV. TIMELINE

Table listing different projects, their timeline and targeted conferences

V. CONCLUSION

Summarize the whole proposal

ACKNOWLEDGMENT